

STATS 3DA3

Homework Assignment 6

Yuhao Wang (400341188)

2024-04-04

Q1:

```
from ucimlrepo import fetch_ucirepo

# fetch dataset
chronic_kidney_disease = fetch_ucirepo(id=336)

# data (as pandas dataframes)
X = chronic_kidney_disease.data.features
y = chronic_kidney_disease.data.targets

# metadata
print(chronic_kidney_disease.metadata)

# variable information
print(chronic_kidney_disease.variables)
```

```
{'uci_id': 336, 'name': 'Chronic Kidney Disease', 'repository_url': 'https://archive.ics.uci.edu/dataset/336/chronic_kidney_disease'}
```

	name	role	type	demographic	description \
0	age	Feature	Integer	Age	None
1	bp	Feature	Integer	None	blood pressure
2	sg	Feature	Categorical	None	specific gravity
3	al	Feature	Categorical	None	albumin
4	su	Feature	Categorical	None	sugar
5	rbc	Feature	Binary	None	red blood cells
6	pc	Feature	Binary	None	pus cell
7	pcc	Feature	Binary	None	pus cell clumps
8	ba	Feature	Binary	None	bacteria
9	bgr	Feature	Integer	None	blood glucose random
10	bu	Feature	Integer	None	blood urea
11	sc	Feature	Continuous	None	serum creatinine
12	sod	Feature	Integer	None	sodium

13	pot	Feature	Continuous	None	potassium
14	hemo	Feature	Continuous	None	hemoglobin
15	pcv	Feature	Integer	None	packed cell volume
16	wbcc	Feature	Integer	None	white blood cell count
17	rbcc	Feature	Continuous	None	red blood cell count
18	htn	Feature	Binary	None	hypertension
19	dm	Feature	Binary	None	diabetes mellitus
20	cad	Feature	Binary	None	coronary artery disease
21	appet	Feature	Binary	None	appetite
22	pe	Feature	Binary	None	pedal edema
23	ane	Feature	Binary	None	anemia
24	class	Target	Binary	None	ckd or not ckd

	units	missing_values
0	year	yes
1	mm/Hg	yes
2	None	yes
3	None	yes
4	None	yes
5	None	yes
6	None	yes
7	None	yes
8	None	yes
9	mgs/dl	yes
10	mgs/dl	yes
11	mgs/dl	yes
12	mEq/L	yes
13	mEq/L	yes
14	gms	yes
15	None	yes
16	cells/cmm	yes
17	millions/cmm	yes

18	None	yes
19	None	yes
20	None	yes
21	None	yes
22	None	yes
23	None	yes
24	None	no

The classification problem using the dataset is to predict whether a patient has chronic kidney disease (CKD) based on 24 medical attributes including age, blood pressure, blood glucose, and more. The target variable is “class” which indicates if the patient has CKD (“ckd”) or not (“notckd”).

Q2:

```
from ucimlrepo import fetch_ucirepo
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier

chronic_kidney_disease = fetch_ucirepo(id=336)

X = chronic_kidney_disease.data.features
y = chronic_kidney_disease.data.targets

numeric_features = [col for col, dtype in zip(X.columns, X.dtypes) if dtype in ['int64', 'float64']]
categorical_features = [col for col in X.columns if col not in numeric_features]

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```

```

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', RandomForestClassifier(random_state=42))])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_test, y_test)
print("Model training complete.")
print("Model accuracy on test set: {:.2f}%".format(accuracy * 100))

```

Model training complete.

Model accuracy on test set: 100.00%

```

d:\python\Lib\site-packages\sklearn\base.py:1474: DataConversionWarning: A column-vector y was
    return fit_method(estimator, *args, **kwargs)

```

Q3:

Variables and Data Types: The dataset consists of 24 features and one target variable “class”: Categorical features: Include specific gravity (sg), albumin (al), sugar (su), red blood cells (rbc), pus cell (pc), pus cell clumps (pcc), so on and The target variable, class, categorizes individuals into

“ckd” (chronic kidney disease) or “notckd” groups. The set has 400 cases, 250 of which have label ‘ckd’, and other 150 which are tagged ‘notckd’, which may require model operators to slightly adjust the algorithms and strategies for training in order to eliminate the bias related to imbalanced class. Dissemination of categorical features likely such “hypertension” and “diabetes mellitus”, which have higher prevalence within the “ckd” category, have probably been referenced to show their association with kidney health.

Q4:

Serum creatinine (sCreat) and blood urea (BUN) are not only considered waste products washed out from the blood by the kidneys in the process of clearance. If these variables show high correlation, elimination of one might not be a matter to discuss because the other variable can cover its lack. Diabetes mellitus (dm), which is a known risk factor for chronic kidney disease (ckd), may significantly influence levels of blood glucose random (bmr) and hemoglobine (hemed). Building interaction terms between ‘dm’ and ‘bgr’, akala ko, ‘dm’ and ‘hemo’ may be a sign of considering the total effect of diabetes on these blood variables. This would likely enhance the model’s operations to forecast CKD in those who are already diabetic. Similarly, htn may often accompany CKD and causes elevated bp levels and sometimes could lead to chronic kidney diseases as indicated by the increased serum creatinine levels. The study of links between hypertension, blood pressure and serum creatinine could be helpful and maybe it could be algorithimized as a predictive factor of chronic kidney disease (CKD).

Q5:

```
from ucimlrepo import fetch_ucirepo

# Check for missing values in the features
missing_values_count = X.isnull().sum()
print("Missing values in each feature:\n", missing_values_count)
```

Missing values in each feature:

age	9
bp	12
sg	47
al	46
su	49
rbc	152
pc	65
pcc	4
ba	4
bgr	44
bu	19
sc	17
sod	87
pot	88
hemo	52
pcv	71
wbcc	106
rbcc	131
htn	2
dm	2
cad	2
appet	1
pe	1
ane	1

dtype: int64

```
from sklearn.impute import SimpleImputer

# Define imputers
# Numeric imputer - using median to avoid influence of outliers
numeric_imputer = SimpleImputer(strategy='median')

# Categorical imputer - using most frequent as it's a common approach for categorical data
```

```

categorical_imputer = SimpleImputer(strategy='most_frequent')

# Impute numerical columns
numerical_columns = X.select_dtypes(include=['int64', 'float64']).columns
X[numerical_columns] = numeric_imputer.fit_transform(X[numerical_columns])

# Impute categorical columns
categorical_columns = X.select_dtypes(include='object').columns
X[categorical_columns] = categorical_imputer.fit_transform(X[categorical_columns])

# Check if any missing values remain
new_missing_values_count = X.isnull().sum()
print("Missing values after imputation:\n", new_missing_values_count)

```

Missing values after imputation:

age	0
bp	0
sg	0
al	0
su	0
rbc	0
pc	0
pcc	0
ba	0
bgr	0
bu	0
sc	0
sod	0
pot	0
hemo	0
pcv	0
wbcc	0


```
rbcc      0
htn       0
dm        0
cad       0
appet     0
pe        0
ane       0
dtype: int64
```

C:\Users\11831\AppData\Local\Temp\ipykernel_15132\95456835.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/

```
X[numerical_columns] = numeric_imputer.fit_transform(X[numerical_columns])
```

C:\Users\11831\AppData\Local\Temp\ipykernel_15132\95456835.py:15: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/

```
X[categorical_columns] = categorical_imputer.fit_transform(X[categorical_columns])
```

Q6:

```
import pandas as pd

X = X.apply(pd.to_numeric, errors='coerce')

# Now check again for missing values after conversion (these might increase)
missing_after_conversion = X.isnull().sum()
print("Missing values after conversion attempt:\n", missing_after_conversion)
```

Missing values after conversion attempt:

age	0
bp	0
sg	0
al	0
su	0
rbc	400
pc	400
pcc	400
ba	400
bgr	0
bu	0
sc	0
sod	0
pot	0
hemo	0
pcv	0
wbcc	0
rbcc	0
htn	400
dm	400
cad	400
appet	400
pe	400
ane	400

dtype: int64

```
# Calculate the IQR (Interquartile Range) to identify outliers
Q1 = X.quantile(0.25)
Q3 = X.quantile(0.75)
IQR = Q3 - Q1
```

```

# Define outliers as those outside of 1.5 * IQR from the Q1 and Q3
outliers = (X < (Q1 - 1.5 * IQR)) | (X > (Q3 + 1.5 * IQR))

# Count outliers in each column
outlier_counts = outliers.sum()
print("Outlier counts in each column:\n", outlier_counts)

# Cap outliers using percentiles
percentiles = X.quantile([0.01, 0.99])
X_capped = X.clip(lower=percentiles.loc[0.01], upper=percentiles.loc[0.99], axis=1)

```

Outlier counts in each column:

age	10
bp	36
sg	7
al	0
su	61
rbc	0
pc	0
pcc	0
ba	0
bgr	53
bu	41
sc	53
sod	18
pot	14
hemo	2
pcv	6
wbcc	17
rbcc	75
htn	0
dm	0

```
cad      0
appet    0
pe        0
ane       0
dtype: int64
```

```
print("Data after capping outliers:\n", X_capped.describe())
# Summary statistics before and after capping
```

Data after capping outliers:

	age	bp	sg	al	su	rbc	pc	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	0.0	0.0	
mean	51.537600	76.300000	1.017712	0.897500	0.387500	NaN	NaN	
std	16.864915	12.193511	0.005434	1.306239	1.009898	NaN	NaN	
min	5.000000	50.000000	1.005000	0.000000	0.000000	NaN	NaN	
25%	42.000000	70.000000	1.015000	0.000000	0.000000	NaN	NaN	
50%	55.000000	80.000000	1.020000	0.000000	0.000000	NaN	NaN	
75%	64.000000	80.000000	1.020000	2.000000	0.000000	NaN	NaN	
max	80.010000	110.000000	1.025000	4.000000	4.000000	NaN	NaN	

	pcc	ba	bgr	...	hemo	pcv	wbcc	\
count	0.0	0.0	400.000000	...	400.000000	400.000000	400.000000	
mean	NaN	NaN	144.709700	...	12.55124	39.107500	8260.690000	
std	NaN	NaN	73.135783	...	2.68300	8.037079	2273.968995	
min	NaN	NaN	70.000000	...	5.79800	16.990000	4097.000000	
25%	NaN	NaN	101.000000	...	10.87500	34.000000	6975.000000	
50%	NaN	NaN	121.000000	...	12.65000	40.000000	8000.000000	
75%	NaN	NaN	150.000000	...	14.62500	44.000000	9400.000000	
max	NaN	NaN	425.220000	...	17.60100	53.010000	16722.000000	

	rbcc	htn	dm	cad	appet	pe	ane
count	400.000000	0.0	0.0	0.0	0.0	0.0	0.0

mean	4.736740	NaN	NaN	NaN	NaN	NaN	NaN
std	0.822137	NaN	NaN	NaN	NaN	NaN	NaN
min	2.499000	NaN	NaN	NaN	NaN	NaN	NaN
25%	4.500000	NaN	NaN	NaN	NaN	NaN	NaN
50%	4.800000	NaN	NaN	NaN	NaN	NaN	NaN
75%	5.100000	NaN	NaN	NaN	NaN	NaN	NaN
max	6.500000	NaN	NaN	NaN	NaN	NaN	NaN

[8 rows x 24 columns]

```
print("Summary statistics before capping:\n", X.describe())
print("Summary statistics after capping:\n", X_capped.describe())
```

Summary statistics before capping:

	age	bp	sg	al	su	rbc	pc	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	0.0	0.0	
mean	51.562500	76.575000	1.017712	0.900000	0.395000	NaN	NaN	
std	16.982996	13.489785	0.005434	1.31313	1.040038	NaN	NaN	
min	2.000000	50.000000	1.005000	0.000000	0.000000	NaN	NaN	
25%	42.000000	70.000000	1.015000	0.000000	0.000000	NaN	NaN	
50%	55.000000	80.000000	1.020000	0.000000	0.000000	NaN	NaN	
75%	64.000000	80.000000	1.020000	2.000000	0.000000	NaN	NaN	
max	90.000000	180.000000	1.025000	5.000000	5.000000	NaN	NaN	

	pcc	ba	bgr	...	hemo	pcv	wbcc	\
count	0.0	0.0	400.000000	...	400.000000	400.000000	400.000000	
mean	NaN	NaN	145.062500	...	12.54250	39.082500	8298.500000	
std	NaN	NaN	75.260774	...	2.71649	8.162245	2529.593814	
min	NaN	NaN	22.000000	...	3.10000	9.000000	2200.000000	
25%	NaN	NaN	101.000000	...	10.87500	34.000000	6975.000000	
50%	NaN	NaN	121.000000	...	12.65000	40.000000	8000.000000	
75%	NaN	NaN	150.000000	...	14.62500	44.000000	9400.000000	

max	NaN	NaN	490.000000	...	17.80000	54.000000	26400.000000
-----	-----	-----	------------	-----	----------	-----------	--------------

		rbcc	htn	dm	cad	appet	pe	ane
count	400.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0
mean	4.737750	NaN	NaN	NaN	NaN	NaN	NaN	NaN
std	0.841439	NaN	NaN	NaN	NaN	NaN	NaN	NaN
min	2.100000	NaN	NaN	NaN	NaN	NaN	NaN	NaN
25%	4.500000	NaN	NaN	NaN	NaN	NaN	NaN	NaN
50%	4.800000	NaN	NaN	NaN	NaN	NaN	NaN	NaN
75%	5.100000	NaN	NaN	NaN	NaN	NaN	NaN	NaN
max	8.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN

[8 rows x 24 columns]

Summary statistics after capping:

		age	bp	sg	al	su	rbc	pc	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	0.0	0.0		
mean	51.537600	76.300000	1.017712	0.897500	0.387500	NaN	NaN		
std	16.864915	12.193511	0.005434	1.306239	1.009898	NaN	NaN		
min	5.000000	50.000000	1.005000	0.000000	0.000000	NaN	NaN		
25%	42.000000	70.000000	1.015000	0.000000	0.000000	NaN	NaN		
50%	55.000000	80.000000	1.020000	0.000000	0.000000	NaN	NaN		
75%	64.000000	80.000000	1.020000	2.000000	0.000000	NaN	NaN		
max	80.010000	110.000000	1.025000	4.000000	4.000000	NaN	NaN		

	pcc	ba	bgr	...	hemo	pcv	wbcc	\
count	0.0	0.0	400.000000	...	400.000000	400.000000	400.000000	
mean	NaN	NaN	144.709700	...	12.55124	39.107500	8260.690000	
std	NaN	NaN	73.135783	...	2.68300	8.037079	2273.968995	
min	NaN	NaN	70.000000	...	5.79800	16.990000	4097.000000	
25%	NaN	NaN	101.000000	...	10.87500	34.000000	6975.000000	
50%	NaN	NaN	121.000000	...	12.65000	40.000000	8000.000000	
75%	NaN	NaN	150.000000	...	14.62500	44.000000	9400.000000	

max	NaN	NaN	425.220000	...	17.60100	53.010000	16722.000000
-----	-----	-----	------------	-----	----------	-----------	--------------

	rbcc	htn	dm	cad	appet	pe	ane
count	400.000000	0.0	0.0	0.0	0.0	0.0	0.0
mean	4.736740	NaN	NaN	NaN	NaN	NaN	NaN
std	0.822137	NaN	NaN	NaN	NaN	NaN	NaN
min	2.499000	NaN	NaN	NaN	NaN	NaN	NaN
25%	4.500000	NaN	NaN	NaN	NaN	NaN	NaN
50%	4.800000	NaN	NaN	NaN	NaN	NaN	NaN
75%	5.100000	NaN	NaN	NaN	NaN	NaN	NaN
max	6.500000	NaN	NaN	NaN	NaN	NaN	NaN

[8 rows x 24 columns]

Q7:

```
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

# Identify numerical and categorical columns
numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns
categorical_cols = X.select_dtypes(include=['object', 'category']).columns

# Create pipelines for the different column types
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
```

```

        ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
        ('onehot', OneHotEncoder(handle_unknown='ignore'))
    ])

# Combine transformers into a single preprocessor object
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)
    ])

# Fit and transform the data
X_preprocessed = preprocessor.fit_transform(X)

```

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without
 warnings.warn(

```

from sklearn.cluster import KMeans

# Determine the optimal number of clusters using the Elbow method
sse = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_preprocessed)
    sse.append(kmeans.inertia_)

# Plotting the Elbow curve
import matplotlib.pyplot as plt


plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), sse, marker='o')
plt.title('Elbow Method')

```



```
plt.xlabel('Number of Clusters')
plt.ylabel('SSE')
plt.show()

# Assuming the elbow is at k = 3
kmeans = KMeans(n_clusters=3, random_state=42)
clusters = kmeans.fit_predict(X_preprocessed)
```



```
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import numpy as np

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_preprocessed)


# Setting up the plot
```

```
plt.figure(figsize=(10, 6))

colors = ['blue', 'green', 'red'] # Define a list of colors for the clusters

# Scatter plot
for i in range(np.max(clusters) + 1): # Loop through the number of clusters
    plt.scatter(X_pca[clusters == i, 0], X_pca[clusters == i, 1],
                color=colors[i], label=f'Cluster {i}', alpha=0.5)

plt.title('Visualization of Patient Clusters')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(title='Cluster') # Add a legend to identify clusters
plt.show()
```



Q8:

```
X_train, X_test = train_test_split(X, test_size=0.30, random_state=1)
print("Training set size:", X_train.shape)
print("Testing set size:", X_test.shape)
```

Training set size: (280, 24)

Testing set size: (120, 24)

Q9: 1,Logistic Regression Logistic Regression provides clear interpretability, which is crucial in medical settings where understanding the influence of predictors is as important as the prediction itself. 2,Random Forest Random Forest can capture complex interactions between features without requiring feature engineering, making it powerful for medical datasets where interactions between symptoms and biological markers can be non-linear and complex.

Q10: Accuracy:it measures the proportion of true results (both true positives and true negatives) among the total number of cases which examined. It provides a indicator of a model's effectiveness at classifying different cases. Confusion Matrices:It helps in understanding exactly where the classifier is making errors, which is critical for medical applications. Knowing the numbers of false positives and false negatives can be crucial for improving diagnostic procedures and treatments.

Q11

```
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=1)

# Identify categorical and numerical columns
categorical_cols = X.select_dtypes(include=['object', 'category']).columns
```

```

numerical_cols = X.select_dtypes(exclude=['object', 'category']).columns

# Preprocessing for numerical data: imputation followed by scaling
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

# Preprocessing for categorical data: imputation followed by one-hot encoding
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)])

model = RandomForestClassifier(n_estimators=100, random_state=42)

pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                            ('model', model)])

# Train the Random Forest Classifier
pipeline.fit(X_train, y_train)

# Predict using the test set
y_pred = pipeline.predict(X_test)

print(y_pred)

```

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without

```

warnings.warn(
d:\python\Lib\site-packages\sklearn\base.py:1474: DataConversionWarning: A column-vector y was
return fit_method(estimator, *args, **kwargs)
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(

['notckd' 'ckd' 'notckd' 'notckd' 'ckd' 'notckd' 'ckd' 'notckd' 'ckd'
'notckd' 'ckd' 'notckd' 'ckd' 'notckd' 'notckd' 'ckd' 'notckd' 'ckd'
'ckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd'
'notckd' 'ckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd' 'ckd' 'ckd' 'ckd'
'notckd' 'ckd' 'notckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd' 'notckd'
'notckd' 'notckd' 'notckd' 'ckd' 'ckd' 'ckd' 'ckd' 'notckd' 'notckd'
'notckd' 'ckd' 'notckd' 'notckd' 'ckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd'
'ckd' 'notckd' 'notckd' 'notckd' 'ckd' 'ckd' 'notckd' 'notckd' 'ckd'
'notckd' 'ckd' 'ckd' 'ckd' 'ckd' 'ckd' 'ckd' 'notckd' 'ckd' 'ckd' 'ckd'
'notckd' 'ckd' 'notckd' 'ckd' 'ckd' 'notckd' 'notckd' 'notckd' 'ckd'
'notckd' 'notckd' 'notckd' 'notckd' 'ckd' 'ckd' 'notckd' 'ckd' 'ckd'
'ckd' 'notckd' 'ckd' 'ckd' 'notckd' 'notckd' 'ckd' 'notckd' 'ckd' 'ckd'
'ckd' 'ckd' 'notckd' 'ckd' 'ckd' 'notckd']

```

```

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

categorical_cols = X.select_dtypes(include=['object', 'category']).columns
numerical_cols = X.select_dtypes(exclude=['object', 'category']).columns

```

```

# Preprocessing for numerical data: imputation followed by scaling
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

# Preprocessing for categorical data: imputation followed by one-hot encoding
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

# Bundle preprocessing for numerical and categorical data
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)])

# Define the models
forest_model = RandomForestClassifier(n_estimators=100, random_state=42)
logreg_model = LogisticRegression(max_iter=1000, random_state=42)

# Create preprocessing and modelling pipelines
forest_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                                   ('model', forest_model)])
logreg_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                                   ('model', logreg_model)])

# Train the Random Forest Classifier
forest_pipeline.fit(X_train, y_train)
# Train the Logistic Regression
logreg_pipeline.fit(X_train, y_train)

```

```
# Predict using the test set with RandomForest
y_pred_forest = forest_pipeline.predict(X_test)
# Predict using the test set with Logistic Regression
y_pred_logreg = logreg_pipeline.predict(X_test)
```

```
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(
d:\python\Lib\site-packages\sklearn\base.py:1474: DataConversionWarning: A column-vector y was
return fit_method(estimator, *args, **kwargs)
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(
d:\python\Lib\site-packages\sklearn\utils\validation.py:1300: DataConversionWarning: A column-
y = column_or_1d(y, warn=True)
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(
```

Q12

```
from sklearn.metrics import accuracy_score, confusion_matrix

# Accuracy for RandomForest
accuracy_forest = accuracy_score(y_test, y_pred_forest)
conf_matrix_forest = confusion_matrix(y_test, y_pred_forest)

# Accuracy for Logistic Regression
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
conf_matrix_logreg = confusion_matrix(y_test, y_pred_logreg)
```

```
# Display the results
print("Random Forest Accuracy:", accuracy_forest)
print("Random Forest Confusion Matrix:\n", conf_matrix_forest)
print("Logistic Regression Accuracy:", accuracy_logreg)
print("Logistic Regression Confusion Matrix:\n", conf_matrix_logreg)
```

Random Forest Accuracy: 1.0

Random Forest Confusion Matrix:

```
[[70  0]
 [ 0 50]]
```

Logistic Regression Accuracy: 0.975

Logistic Regression Confusion Matrix:

```
[[67  3]
 [ 0 50]]
```

Q13

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline

# Define which columns are categorical and numerical in your dataset
categorical_cols = X.select_dtypes(include=['object', 'bool', 'category']).columns
numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns

# Preprocessing for numerical data: imputation followed by scaling
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```



```
# Preprocessing for categorical data: imputation followed by one-hot encoding
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

# Bundle preprocessing for numerical and categorical data
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)])

# Apply transformations to the data
X_preprocessed = preprocessor.fit_transform(X)
```

```
d:\python\Lib\site-packages\sklearn\impute\_base.py:577: UserWarning: Skipping features without
warnings.warn(
```

```
from sklearn.linear_model import LogisticRegression

# Initialize the Logistic Regression model
logreg_full = LogisticRegression(random_state=1)

# Fit the model on the preprocessed data
logreg_full.fit(X_preprocessed, y)
```

```
d:\python\Lib\site-packages\sklearn\utils\validation.py:1300: DataConversionWarning: A column-
y = column_or_1d(y, warn=True)
```

```
LogisticRegression(random_state=1)
```

```

# Get the coefficients from the trained model
feature_names = preprocessor.get_feature_names_out() # Get the feature names after preprocessing
coefficients = logreg_full.coef_[0]

# Associate each coefficient with its corresponding feature name
importance = pd.Series(coefficients, index=feature_names)

# Sort features by their coefficient magnitude for interpretation
sorted_importance = importance.abs().sort_values(ascending=False)
print(sorted_importance)

```

```

num__hemo    0.987752
num__al      0.925141
num__sg      0.882522
num__sc      0.714171
num__bu      0.659715
num__su      0.625505
num__pcv     0.519588
num__bgr     0.333517
num__bp      0.270300
num__age     0.244780
num__sod     0.235290
num__pot     0.127910
num__wbcc    0.091388
num__rbcc    0.026597
dtype: float64

```

The features with the highest coefficients are generally those directly related to kidney function tests and blood tests, which are critical in diagnosing and managing CKD. This model effectively highlights the key biomarkers for CKD, which could help in early detection and management strategies.

Q14

As for interaction terms or the polynomial features are concerned they can be employed to ensure the non-linear relationships between them are taken into account. Finally, after the application of learning mechanisms, apply accuracy and confusion matrices to evaluate the upgrade process. Compare these results to the previously used models for the targeted sub-groups to check whether the performance of the model parametrically upgraded or not.

```
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

categorical_cols = [col for col in X.columns if X[col].dtype == 'object']
numerical_cols = [col for col in X.columns if X[col].dtype != 'object'] # Numerical columns

# Create transformers for preprocessing
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Use median for numerical columns
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Use 'missing' for
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

# Bundle preprocessing for numerical and categorical data
preprocessor = ColumnTransformer(
```

```

transformers=[
    ('num', numeric_transformer, numerical_cols),
    ('cat', categorical_transformer, categorical_cols)
])

# Define the classifiers
random_forest_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
logistic_regression_classifier = LogisticRegression(max_iter=1000, random_state=42)

# Create preprocessing and modeling pipelines
pipeline_rf = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier', random_forest_classifier)])

pipeline_lr = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier', logistic_regression_classifier)])

# Fit both classifiers
pipeline_rf.fit(X_train, y_train)
pipeline_lr.fit(X_train, y_train)

# Predict using both classifiers
y_pred_rf = pipeline_rf.predict(X_test)
y_pred_lr = pipeline_lr.predict(X_test)

# Evaluate both classifiers
accuracy_rf = accuracy_score(y_test, y_pred_rf)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

accuracy_lr = accuracy_score(y_test, y_pred_lr)
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

```

```

print(f"Random Forest Accuracy: {accuracy_rf}")
print("Random Forest Confusion Matrix:\n", conf_matrix_rf)

print(f"Logistic Regression Accuracy: {accuracy_lr}")
print("Logistic Regression Confusion Matrix:\n", conf_matrix_lr)

```

Random Forest Accuracy: 1.0

Random Forest Confusion Matrix:

```

[[70  0]
 [ 0 50]]

```

Logistic Regression Accuracy: 0.975

Logistic Regression Confusion Matrix:

```

[[67  3]
 [ 0 50]]

```

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without
warnings.warn(

d:\python\Lib\site-packages\sklearn\base.py:1474: DataConversionWarning: A column-vector y was
return fit_method(estimator, *args, **kwargs)

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without
warnings.warn(

d:\python\Lib\site-packages\sklearn\utils\validation.py:1300: DataConversionWarning: A column-
y = column_or_1d(y, warn=True)

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without
warnings.warn(

d:\python\Lib\site-packages\sklearn\impute_base.py:577: UserWarning: Skipping features without
warnings.warn(

Grading scheme

1. Answer [1]
2. Codes [2]
OR answer [2]
3. Codes [3] and answer [3]
4. Codes [2] and answer [3]
5. Codes [2]
OR answer [2]
6. Codes [2]
OR answer [2]
7. Codes [3] and Plot [1]
8. Codes [1]
9. Answers [2]
10. Describe the two metrics [2]
11. Codes [2]
these codes can be included in (12)
12. Codes (two classifiers training,
model selection for each classifier,
classifiers comparisons) [5] and answer [2]
13. Codes [1] and answers [2]
14. Codes and comparison will
give **bonus 2 points for the final grade.**

The maximum point for this assignment is 39. We will convert this to 100%.

All group members will receive the same grade if they contribute to the same.